

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Stack frames, how hard can it be?

dram

2023-04-19

About me

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- @dramforever on most random platforms
- <https://dram.page>
- Call me 'dram'

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Basics

Calling convention

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

At function entry:

- ra (x1) return address
- sp (x2) stack pointer

At function exit:

- Restore sp
- Return value in a0
- Jump to original ra

Stack grows down

Just a random function

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

```
long answer(void) {  
    return 42;  
}
```

Just a random function

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

```
.section .text
```

```
.global answer
```

```
.type answer, @function
```

```
answer:
```

```
li a0, 42
```

```
ret
```

```
.size answer, . - answer
```

Leaf functions

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Leaf functions...

- Don't have to use the stack
- (But can if they need to)
- No need to spill ra on the stack

More complex functions

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

```
int foo();  
  
int bar() {  
    return foo() + foo();  
}
```


More complex functions

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

bar:

```
addi sp, sp, -16
```

```
sd ra, 8(sp)
```

```
sd s0, 0(sp)
```

```
call foo
```

```
mv s0, a0           # s0 = foo()
```

```
call foo
```

```
addw a0, a0, s0     # a0 = s0 + foo()
```

```
ld ra, 8(sp)
```

```
ld s0, 0(sp)
```

```
addi sp, sp, 16
```

```
ret
```

More complex functions

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

At function entry (i.e. prologue):

- Decrement sp
- Save ra and used callee-saved regs on stack

At function exit (i.e. epilogue):

- Restore ra and callee-saved regs from stack
- Increment sp
- Return to ra

Wait, is ra callee-saved?

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- Calling convention spec says no! ra is not preserved across calls
- `jal ra, ...` and `jalr ra, ...` destroys ra anyway

Wait, is ra callee-saved?

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Technically allowed

```
ld t2, 8(sp)    # Load return address
jr t2
```

- Save ra anyway for non-leaf functions, because we need to know where to return to
- Restore ra anyway for return, because jr ra is recognized as a return for RAS prediction

Bonus: What do the insns look like?

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- No stack usage, no stack frame allocation needed
- Small stack frame (< 512 bytes): `c.addi16sp`
- Larger frame (< 2048 bytes): `addi`

Bonus: What do the insns look like?

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- No stack usage, no stack frame allocation needed
- Small stack frame (< 512 bytes): `c.addi16sp`
- Larger frame (< 2048 bytes): `addi`
- Very large frame (< 2 GiB):

```
li t0, -1000000 # lui/addi
sub sp, sp, t0
```

Bonus: What do the insns look like?

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- Large-ish frames (< ~ 6000 bytes?):

```
# LLVM
addi    sp, sp, -2032
sd      ra, 2024(sp)
addi    sp, sp, -2048
addi    sp, sp, -944
```

If advantageous, may generate multiple addi instructions, bump sp in multiple steps

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Frame pointer

Frame pointer

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

On RISC-V, `x8/s0/fp` is the frame pointer.

Standardized in April 2023 (!), frame pointer convention¹:

- After function prologue, `fp` equals the CFA (canonical frame address), which is `sp` at function entry
- RV64: Previous `fp` at `-16(fp)`, saved `ra` at `-8(fp)`
- (RV32: Previous `fp` at `-8(fp)`, saved `ra` at `-4(fp)`)

```
addi sp, sp, - N      # c.addi16sp
sd ra, (N - 8)(sp)    # c.sdsp
sd s0, (N - 16)(sp)   # c.sdsp
addi s0, sp, N        # c.addi4spn
```

¹<https://github.com/riscv-non-isa/riscv-elf-psabi-doc/pull/369>

Frame pointer

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Both GCC and LLVM force a frame pointer if the function uses `alloca` or VLA:

```
void bar(size_t n) {  
    char buf[n]; // Variable-length array  
    foo(buf);  
}
```

Force enabling fp

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- gcc and clang: `-fno-omit-frame-pointer`
- rustc: `-Cforce-frame-pointers=yes`

Differences in GCC and LLVM

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

For leaf functions, GCC (as of 12.2.0) does not save `s0` and `ra` in the ‘correct’ place even with `-fno-omit-frame-pointer`². LLVM does.

(They both do correctly have `fp = original sp`.)

²<https://godbolt.org/z/vK7TTqTae>

Differences in GCC and LLVM

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

```
# GCC 12.2.0
```

```
answer:
```

```
addi sp, sp, -16
```

```
sd    s0, 8(sp)
```

```
addi s0, sp, 16
```

```
# ...
```

```
# Clang 16.0.0
```

```
answer:
```

```
addi sp, sp, -16
```

```
sd    ra, 8(sp)
```

```
sd    s0, 0(sp)
```

```
addi s0, sp, 16
```

```
# ...
```

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Code size optimization

Code size optimization

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Prologues and epilogues take up code size

- `__riscv_{save,restore}_N` out-of-line prologue and epilogue
- Zcmp extension

__riscv_{save,restore}_N

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

```
int bar() { return foo() + foo(); }
```

gcc -O -msave-restore gives:

```
bar:  
    call    t0, __riscv_save_1  
    call    foo  
    mv      s0, a0  
    call    foo  
    addw    a0, s0, a0  
    tail    __riscv_restore_1
```


__riscv_{save,restore}_N

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- __riscv_save_N: Push the first N callee-saved registers and ra, returns to t0
- __riscv_restore_N: Pop the first N callee-saved registers and ra, returns to ra

__riscv_save_N uses the alternate link register t0 so calling these does not have to destroy ra:

```
__riscv_save_1:  
    addi sp, sp, -16  
    sd s0, 0(sp)  
    sd ra, 8(sp)  
    jr t0
```

Zcmp extension

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

The Zcmp extension³ defines `cm.push`, `cm.pop` etc

```
bar:
    cm.push    {ra, s0}, -16
    # ...
    cm.popret  {ra, s0}, 16
```

(ARM: `push {lr, ...}` and `pop {pc, ...}`)

³<https://github.com/riscv/riscv-code-size-reduction>

Zcmp extension

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

```
# cm.push {ra, s0}, -16
```

```
sd s0, -8(sp)
```

```
sd ra, -16(sp)
```

```
addi sp, sp, -16
```

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

A small problem

Zcmp stack frame layout

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Zcmp stack frame layout:

```
# cm.push {ra, s0}, -16  
sd s0, -8(sp)  
sd ra, -16(sp)  
addi sp, sp, -16
```

Frame pointer convention and __riscv_{save,restore}_N:

answer:

```
addi sp, sp, -16  
sd ra, 8(sp)  
sd s0, 0(sp)
```

Stack frame layout

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- Zcmp is incompatible with `-fno-omit-frame-pointer` — must use longer sequence in this case
- `-msave-restore -fno-omit-frame-pointer` causes GCC (as of 12.2.0) to ignore `-msave-restore`, while LLVM inserts a fp update after `__riscv_save_N`

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

Conclusion

Conclusion

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- There's lots of freedom in the base RISC-V calling convention about how the stack is structured
- Frame pointer convention essential for fast unwinding (perf, garbage collection, easy stack backtraces ...)
- Unfortunately most reliable unwinding is still DWARF...

References

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion

- Calling convention: <https://github.com/riscv-non-isa/riscv-elf-psabi-doc>
- Zcmp: <https://github.com/riscv/riscv-code-size-reduction>

Thank you

Stack frame

dram

Basics

Frame pointer

Code size
optimization

A small
problem

Conclusion